

## APPENDIX: CODE

```
"""
```

```
EECS 445 - Introduction to Machine Learning
```

```
Winter 2019
```

```
Homework 2, Ensemble Methods
```

```
Skeleton Code
```

```
"""
```

```
import random
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
from collections import Counter
```

```
from sklearn import metrics, utils
```

```
from sklearn.datasets import fetch_mldata
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.model_selection import train_test_split
```

```
def load_mnist(classes):
```

```
    """
```

```
    Load MNIST dataset for classes
```

```
    Every 25th sample is used to reduce computational resources
```

```
    Input:
```

```
        classes : list of ints
```

```
    Returns:
```

```
        X : np.array (num_samples, num_features)
```

```
        y : np.array (num_samples)
```

```
    """
```

```
    print('Fetching MNIST data...')
```

```
    mnist = fetch_mldata('MNIST original')
```

```
    X_all = np.array(mnist.data)[::25]
```

```
    y_all = np.array(mnist.target)[::25]
```

```
    desired_idx = np.isin(y_all, classes)
```

```
    return X_all[desired_idx], y_all[desired_idx]
```

```
def get_avg_performance(X, y, m_vals, n_splits=50):
```

```
    """
```

```
    Compare the average performance of bagging and random forest across 50
```

```
    random splits of X and y
```

```
    Input:
```

```
        X : np.array (num_samples, num_features)
```

```
        y : np.array (num_samples)
```

```
        m_vals: list - list of values for m
```

```
        n_splits: int - number of random splits
```

```
    Returns:
```

```
        bag_results : np.array (len(m_vals)) - estimate of bagging performance
```

```
        rf_results : np.array (len(m_vals)) - estimate of random forest performance
```

```
    """
```

```
    print('Getting bagging and random forest scores...')
```

```
    rf_results = []
```

```
    bag_results = []
```

```
    for m in m_vals:
```

```
        print('m = {}'.format(m))
```

```

bagging_scores = []
random_forest_scores = []
for i in range(n_splits):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
    random_forest_scores.append(random_forest(X_train, y_train, X_test, y_test, m))
    bagging_scores.append(bagging_ensemble(X_train, y_train, X_test, y_test))
rf_results.append(np.median(np.array(random_forest_scores)))
bag_results.append(np.median(np.array(bagging_scores)))
return bag_results, rf_results

```

```

def plot_data(bagging_scores, random_forest_scores, m_vals):

```

```

    """

```

Plot bagging and random forest accuracies

Input:

bagging\_scores : np.array - array containing accuracies for bagging ensemble classifiers

random\_forest\_scores : np.array - array containing accuracies for random forest classifiers

```

    """

```

```

plt.figure()
plt.plot(list(m_vals), bagging_scores, '--', label='bagging')
plt.plot(list(m_vals), random_forest_scores, '--', label='random forest')
plt.xlabel('m')
plt.ylabel('Accuracy')
plt.legend(loc='upper right')
plt.savefig('ensemble.png', dpi=300)
plt.show()

```

```

def random_forest(X_train, y_train, X_test, y_test, m, n_clf=10):

```

```

    """

```

Returns accuracy on the test set X\_test with corresponding labels y\_test using a random forest classifier with n\_clf decision trees trained with training examples X\_train and training labels y\_train.

Input:

X\_train : np.array (n\_train, d) - array of training feature vectors

y\_train : np.array (n\_train) - array of labels corresponding to X\_train samples

X\_test : np.array (n\_test, d) - array of testing feature vectors

y\_test : np.array (n\_test) - array of labels corresponding to X\_test samples

m : int - number of features to consider when splitting

n\_clf : int - number of decision tree classifiers in the random forest, default is 10

Returns:

accuracy : float - accuracy of random forest classifier on X\_test samples

```

    """

```

# TODO: Implement this function

```

y_predict = np.zeros((10,X_test.shape[0]))
boot_size = X_train.shape[0]
for i in range(10):
    X_boot, y_boot = utils.resample(X_train, y_train, n_samples = boot_size)
    clf = DecisionTreeClassifier(criterion = 'entropy', max_features = m)
    clf.fit(X_boot, y_boot)
    y_pred = clf.predict(X_test)
    y_predict[i] = y_pred

y_pred = []
for i in range(X_test.shape[0]):
    y_pred = np.append(y_pred,Counter(y_predict[:,i]).most_common(1)[0][0])

```

```
return metrics.accuracy_score(y_test, y_pred)
```

```
def bagging_ensemble(X_train, y_train, X_test, y_test, n_clf=10):
```

```
    """
```

```
    Returns accuracy on the test set X_test with corresponding labels y_test
    using a bagging ensemble classifier with n_clf decision trees trained with
    training examples X_train and training labels y_train.
```

```
    Input:
```

```
    X_train : np.array (n_train, d) - array of training feature vectors
    y_train : np.array (n_train) - array of labels corresponding to X_train samples
    X_test : np.array (n_test,d) - array of testing feature vectors
    y_test : np.array (n_test) - array of labels corresponding to X_test samples
    n_clf : int - number of decision tree classifiers in the random forest, default is 10
```

```
    Returns:
```

```
    accuracy : float - accuracy of random forest classifier on X_test samples
```

```
    """
```

```
    # TODO: Implement this function
```

```
    y_predict = np.zeros((10,X_test.shape[0]))
```

```
    boot_size = X_train.shape[0]
```

```
    for i in range(10):
```

```
        X_boot, y_boot = utils.resample(X_train, y_train, n_samples = boot_size)
```

```
        clf = DecisionTreeClassifier(criterion = 'entropy')
```

```
        clf.fit(X_boot, y_boot)
```

```
        y_pred = clf.predict(X_test)
```

```
        y_predict[i] = y_pred
```

```
    y_pred = []
```

```
    for i in range(X_test.shape[0]):
```

```
        y_pred = np.append(y_pred,Counter(y_predict[:,i]).most_common(1)[0][0])
```

```
    return metrics.accuracy_score(y_test, y_pred)
```

```
def main():
```

```
    """
```

```
    Analyze how the performance of bagging and random forest changes with m.
```

```
    """
```

```
    X, y = load_mnist([1,2,3,4])
```

```
    # Plot accuracies
```

```
    m_vals = [1] + list(range(56, 785, 56))
```

```
    bagging_scores, random_forest_scores = get_avg_performance(X, y, m_vals)
```

```
    plot_data(bagging_scores, random_forest_scores, m_vals)
```

```
if __name__ == '__main__':
```

```
    main()
```